

Session 5

How to interpret the latent class

Matt Denwood

2021-06-30

Session 5: How to interpret the latent class

Recap

- Adding more populations and more tests to a Hui-Walter model is technically easy
 - Particularly if using `template_huiwalter`
- Verifying that the assumptions you are making are correct is harder
 - The sensitivity and specificity must be consistent across populations
 - Pairwise correlation between tests should be accounted for (with >2 tests)

What exactly is our latent class?

Homework (reminder): think about what exactly the latent class is in these situations:

1. An antigen plus antibody test

What exactly is our latent class?

Homework (reminder): think about what exactly the latent class is in these situations:

1. An antigen plus antibody test
 - The latent status is probably close to the true disease status

What exactly is our latent class?

Homework (reminder): think about what exactly the latent class is in these situations:

1. An antigen plus antibody test
 - The latent status is probably close to the true disease status
1. Two antibody tests

What exactly is our latent class?

Homework (reminder): think about what exactly the latent class is in these situations:

1. An antigen plus antibody test

- The latent status is probably close to the true disease status

1. Two antibody tests

- The latent status is actually 'producing antibodies' not 'diseased'

What exactly is our latent class?

Homework (reminder): think about what exactly the latent class is in these situations:

1. An antigen plus antibody test

- The latent status is probably close to the true disease status

1. Two antibody tests

- The latent status is actually ‘producing antibodies’ not ‘diseased’
- What do we mean by “conditionally independent” (revisited) ?
 - Independent of each other conditional on the latent state
 - But the latent state is NOT always *disease*

What exactly is our latent class?

Homework (reminder): think about what exactly the latent class is in these situations:

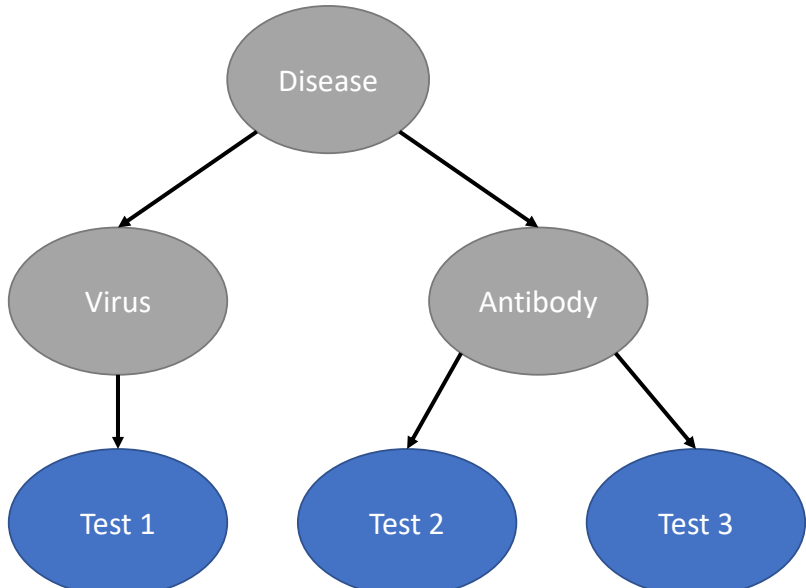
1. An antigen plus antibody test

- The latent status is probably close to the true disease status

1. Two antibody tests

- The latent status is actually ‘producing antibodies’ not ‘diseased’
- What do we mean by “conditionally independent” (revisited) ?
 - Independent of each other conditional on the latent state
 - But the latent state is NOT always *disease*
- i.e. we’re pulling **something** out of a hat, and deciding to call it a rabbit

A hierarchy of latent states



Branching of processes leading to test results

- Sometimes we have multiple tests that are detecting a similar thing
 - For example: two antibody tests and one antigen test
 - The antibody tests will be correlated

Branching of processes leading to test results

- Sometimes we have multiple tests that are detecting a similar thing
 - For example: two antibody tests and one antigen test
 - The antibody tests will be correlated
- Or even three antibody tests where two are primed to detect the same thing, and one has a different target!
 - In this case all three tests are correlated, but two are more strongly correlated

Optional Exercise Code

*# Probability of antibody response conditional on disease status (really
↳ bad to illustrate the point):*

```
se_antibody <- 0.5  
sp_antibody <- 0.75  
N <- 100
```

Otherwise the parameters are as before

True latent infection status as before:

```
true <- rbinom(N, 1, prevalence[Group])
```

Latent class of antibody response conditional on the true status:

```
antibody <- rbinom(N, 1, se_antibody*true + (1-sp_antibody)*(1-true))
```

Simulate test results for test 1 conditional on antibody status:

```
test1 <- rbinom(N, 1, se1*antibody + (1-sp1)*(1-antibody))
```

etc

Note that the overall sensitivity and specificity of the tests needs to

↳ be corrected for the antibody positive step:

```
overall_se1 <- se_antibody*se1 + (1-se_antibody)*(1-sp1)  
overall_sp1 <- sp_antibody*sp1 + (1-sp_antibody)*(1-se1)
```

etc

Optional Solution

```
# Parameter values to simulate:
```

```
N <- 200
```

```
se1 <- 0.8
```

```
sp1 <- 0.95
```

```
se2 <- 0.9
```

```
sp2 <- 0.99
```

```
se3 <- 0.95
```

```
sp3 <- 0.95
```

```
# Probability of antibody response conditional on disease status (really  
↪ bad to illustrate the point):
```

```
se_antibody <- 0.5
```

```
sp_antibody <- 0.75
```

```
Populations <- 2
```

```
prevalence <- c(0.25,0.75)
```

```
Group <- sample(1:Populations, N, replace=TRUE)
```

```
# Ensure replicable data:
```

```
set.seed(2020-02-18)
```

```
# True latent infection status as before:
```

```
true <- rbinom(N, 1, prevalence[Group])
```

What is sensitivity and specificity

TODO: show how overall Se and Sp relates to probability of test positive given antibody response etc

Alternative model formulation

TODO: change to explicit serial Disease -> Antibody -> Test etc

Note that autocorrelation is terrible

```
model{  
  
  for(i in 1:N){  
    Status[i] ~ dcat(prob[i, ])  
  
    prob[i,1] <- (prev[i] * ((1-se[1])*(1-se[2]))) +  
                 ((1-prev[i]) * ((sp[1])*(sp[2])))  
    prob[i,2] <- (prev[i] * ((se[1])*(1-se[2]))) +  
                 ((1-prev[i]) * ((1-sp[1])*(sp[2])))  
    prob[i,3] <- (prev[i] * ((1-se[1])*(se[2]))) +  
                 ((1-prev[i]) * ((sp[1])*(1-sp[2])))  
    prob[i,4] <- (prev[i] * ((se[1])*(se[2]))) +  
                 ((1-prev[i]) * ((1-sp[1])*(1-sp[2])))  
  
    logit(prev[i]) <- intercept + population_effect[Population[i]]  
  }  
}
```

```

intercept ~ dnorm(0, 0.33)
population_effect[1] <- 0
for(p in 2:Pops){
  population_effect[p] ~ dnorm(0, 0.1)
}
se[1] ~ dbeta(1, 1)T(1-sp[1], )
sp[1] ~ dbeta(1, 1)
se[2] ~ dbeta(1, 1)T(1-sp[2], )
sp[2] ~ dbeta(1, 1)

#data# Status, N, Population, Pops
#monitor# intercept, population_effect, se, sp
#inits# intercept, population_effect, se, sp
}

```

- The main difference is the prior for prevalence in each population
- We also need to give initial values for intercept and population_effect rather than prev, and tell run.jags the data frame from which to extract the data (except N and Pops):

```
intercept <- list(chain1=-1, chain2=1)
population_effect <- list(chain1=c(NA, 1, -1), chain2=c(NA, -1, 1))
se <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
sp <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))

simdata$Status <- with(simdata, factor(interaction(Test1, Test2),
  ↳ levels=c('0.0','1.0','0.1','1.1')))
N <- nrow(simdata)
Pops <- length(levels(simdata$Population))
glm_results <- run.jags('glm_hw.bug', n.chains=2, data=simdata)
```

Publication of your results

STARD-BLCM: A helpful structure to ensure that papers contain all necessary information

If you use the software, please cite JAGS:

Plummer, M. (2003). JAGS : A Program for Analysis of Bayesian Graphical Models Using Gibbs Sampling JAGS : Just Another Gibbs Sampler. Proceedings of the 3rd International Workshop on Distributed Statistical Computing (DSC 2003), March 20–22, Vienna, Austria. ISSN 1609-395X. <https://doi.org/10.1.1.13.3406>

R:

```
citation()  
##  
## To cite R in publications use:  
##  
## R Core Team (2021). R: A language and environment  
## for statistical computing. R Foundation for
```

Discussion session 5

Points to consider

1. Interpreting the results of latent class models is much more difficult than running them
2. How can we be sure that e.g. probability of a positive test result conditional on the latent state is the same thing as sensitivity?
3. How can we make sure that our publications contain all of the necessary information to allow others to interpret our findings?

Exercise

1. Read the STARD-BLCM guidelines, checklist, and examples documents
2. Read the Diagnosing diagnostic tests paper
3. Be ready with questions for the group discussion!

Summary

- Latent class models are MUCH more complex to interpret than traditional models
 - Take time to think about what the latent class means
- Think about which tests might be correlated and if you should include covariance terms
- Think about the biology of where your data comes from, particularly if populations are fundamentally different
- Follow the STARD checklist!